

# Introduction to Applied Reinforcement Learning

Shah Rukh Qasim


Department of Mathematical Modeling and Machine Learning  
&&  
Physik Institut  
University of Zurich

[shahrukh.qasim@physik.uzh.ch](mailto:shahrukh.qasim@physik.uzh.ch)

25.11.2025

# Today's Agenda

- We will skip the tutorial today
  - Although I will give a link to an example towards the end to play
  - Interest of time: there is a lot to learn to get started with RL
- Reinforcement Learning (RL) is hard and there is a lot to learn
  - We will cover the core concepts
- The order of the contents of this lecture is non traditional in how RL is taught: Fundamentals are discussed later
  - Expectations is that it will make it easier to understand quickly
- References in Sutton and Barto 2nd Edition
  - Freely available at: <http://incompleteideas.net/book/the-book-2nd.html>




<https://tinyurl.com/AppliedRLNotebook>  
[https://colab.research.google.com/drive/1suKXuHlf6M0N2yh1K\\_OAq-qJwpMe4ytM?usp=sharing](https://colab.research.google.com/drive/1suKXuHlf6M0N2yh1K_OAq-qJwpMe4ytM?usp=sharing)

# Why Reinforcement Learning?

- Playing chess
  - Input: state of the board
  - Output: the next move
  - Supervised learning:
    - Collect a whole bunch of samples with the next best move and then train the NN on it
  - Reinforcement learning:
    - Let the agent make decisions
    - ... to maximize the reward function
    - It's the job of the agent to collect data and take the best decision at every point



# Reinforcement Learning




- RL agent gets a state  $s_t$
- Takes an action  $a_t$
- Gets a reward  $r_{t+1}$
- The state gets updated  $s_{t+1}$

Remember: At every time step, an agent makes a decision – **ONLY** based on the current state!

If the history is important, append it to the state!

# Reinforcement Learning

- Capital letters represent random variables:  
not exact, we blur the boundaries
- An episode is then:
  - $(s_0, a_0, r_1, s_1, a_1, r_2, s_2, \dots, a_{T-1}, r_T, s_T)$
- Maximize:
  - $G_t = R_{t+1} + R_{t+2} + \dots + R_T$ 
    - (incomplete)
  - Cumulative reward
  - Also called return
  - $G_0$  is episode return



Agent's Goal: Maximize Return  $G_t$  at every step, not  $R_{t+1}$

# Discounted Returns

- We need to add a discount factor
  - $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T$
  - $\gamma \leq 1$ , immediate rewards matter more
    - You can invest money now.
    - There is inflation.
    - There is uncertainty about the future.
- Or infinitely long episodes:
  - $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$

# Policy

- What is a policy?
  - A rule or strategy that tells an agent how to act in each situation.
    - State: A customer asks to return an item (with or without receipt)
    - Policy: "If the customer has a receipt, accept the return; otherwise, decline"
- In RL:
  - A policy can be **deterministic**
    - $a_t = \pi(s_t)$
    - $\pi : \mathcal{S} \rightarrow \mathcal{A}$
  - A policy can be **stochastic**
    - $\pi(a | s) = \Pr(A_t = a | S_t = s)$
    - $\pi : \mathcal{S} \times \mathcal{A} \rightarrow [0,1]$

## Why stochastic?

- A deterministic policy can cause a vacuum robot to loop forever by repeating the same action in the same state, while a stochastic policy can break the loop.
- When several actions are equally good, a stochastic policy represents this better than forcing one arbitrary choice.
- Stochasticity enables learning (discussed later)

# Policy Gradient: REINFORCE

- The NN is
  - $\pi_\theta$ : It products output logits – same as a normal classification task
  - It can also be a continuous output
- Sample full episodes
- Compute returns
- Change policies to make actions more likely which led to higher returns
  - Proof: 13.3 in Sutton and Barto
- Wait till the end of episode: **Monte Carlo** learning

---

**Algorithm 1** REINFORCE (Monte Carlo Policy Gradient) with Sampled Trajectories

---

**Require:** learning rate  $\alpha$ , discount factor  $\gamma$ , initial parameters  $\theta$

```
1: for episodes  $n = 1, 2, \dots$  do
2:   Generate a trajectory  $(s_0, a_0, r_1, s_1, \dots, s_T)$  using policy  $\pi_\theta(a_t | s_t)$ :
3:   for  $t = 0$  to  $T - 1$  do
4:     Sample action  $a_t \sim \pi_\theta(\cdot | s_t)$ 
5:     Environment produces reward  $r_{t+1}$  and next state  $s_{t+1}$ 
6:   end for
7:   Compute returns  $G_t$  for  $t = T - 1, \dots, 0$ :
8:   for  $t = T - 1$  down to 0 do
9:      $G_t \leftarrow r_{t+1} + \gamma G_{t+1}$                                      (with  $G_T \leftarrow 0$ )
10:    end for
11:    Compute policy gradient estimate:
12:      
$$g \leftarrow \sum_{t=0}^{T-1} G_t \nabla_\theta \log \pi_\theta(a_t | s_t)$$

13:    Update parameters:
14:      
$$\theta \leftarrow \theta + \alpha g$$

15:  end for
```

---

# Value function

- An agent can predict how promising each situation (state) is in the long run.
  - Done through state-value function:
    - $V^\pi(s) = \mathbb{E}_\pi [ G_t \mid S_t = s ]$
    - The state-value function tells us how good it is to be in a given state when the agent follows a particular policy  $\pi$
    - It is the expected return starting from state  $s$  and acting according to  $\pi$  thereafter
  - Or  $Q$  function
    - $Q^\pi(s, a) = \mathbb{E}_\pi [ G_t \mid S_t = s, A_t = a ]$

# REINFORCE with baseline

- Why does it help?
  - Without a baseline: grading students on raw scores
  - Baseline: grading relative to class average
- Reduces the variance
- This is an **Actor–Critic** algorithm
  - the policy network acts by selecting actions, while the value network critiques those actions by estimating their expected returns and providing the advantage signal for learning
- Unfortunately, REINFORCE is **not** the foundation of modern RL algorithms

---

**Algorithm 2** REINFORCE with Learned Baseline (Monte Carlo Actor–Critic)

**Require:** learning rate  $\alpha$  for policy, learning rate  $\beta$  for value network, discount factor  $\gamma$

**Require:** initial policy parameters  $\theta$ , initial value network parameters  $\phi$

- 1: **for** episodes  $n = 1, 2, \dots$  **do**
- 2:   Generate a trajectory  $(s_0, a_0, r_1, s_1, \dots, s_T)$  using the policy network  $\pi_\theta(a_t | s_t)$  (*and record baseline estimates  $V_\phi(s_t)$  along the way*):
- 3:   **for**  $t = 0$  to  $T - 1$  **do**
- 4:     Sample action  $a_t \sim \pi_\theta(\cdot | s_t)$
- 5:     Evaluate baseline  $V_\phi(s_t)$
- 6:     Environment produces reward  $r_{t+1}$  and next state  $s_{t+1}$
- 7:   **end for**
- 8:   Compute returns  $G_t$  for  $t = T - 1, \dots, 0$ :
- 9:   **for**  $t = T - 1$  down to 0 **do**
- 10:      $G_t \leftarrow r_{t+1} + \gamma G_{t+1}$  (with  $G_T \leftarrow 0$ )
- 11:   **end for**
- 12:   Compute advantages:  
$$A_t \leftarrow G_t - V_\phi(s_t)$$

- 13:   Update the value network:

$$\phi \leftarrow \phi - \beta \sum_{t=0}^{T-1} \nabla_\phi (G_t - V_\phi(s_t))^2$$

- 14:   Update the policy network:


$$\theta \leftarrow \theta + \alpha \sum_{t=0}^{T-1} A_t \nabla_\theta \log \pi_\theta(a_t | s_t)$$

# Value based control

- Control = Simply Learning (a policy or otherwise)
- Instead of policy gradient  $[a_t \sim \pi_\theta(\cdot | s_t)]$ 
  - We can also do value based learning with the  $Q$  function. Here we only have  $Q(s, a)$  defined as a NN and action selection has to be added externally
    - $a_t = \begin{cases} \text{a random action from } \mathcal{A}(s_t), & \text{with probability } \varepsilon, \\ \arg \max_a Q(s_t, a), & \text{with probability } 1 - \varepsilon. \end{cases}$
    - $Q^\pi(s, a) = \mathbb{E}_\pi \left[ G_t \mid S_t = s, A_t = a \right]$
  - Exploration vs exploitation
    - Why we don't have  $\epsilon$  greedy in policy gradient?
    - We can add entropy:

$$H(\pi(\cdot | s)) = - \sum_a \pi(a | s) \log \pi(a | s)$$

Example if you have only two actions



# Temporal Difference Learning

- Foundation of RL: Temporal Difference Learning
- Let's first explain with a real world example:
  - What happens if I move the queen to the red star?
    - Do I need to wait for the game to finish?
  - What happens if I take a wrong turn and end up on the opposite end of the highway?
    - Do I need to wait for an accident to happen?
- Also called **bootstrapping**
- Why is this better?



# SARSA

- See how we are learning at every step instead of waiting for the end?
    - That's temporal difference learning
  - Reading: TD- $\lambda$ :
    - Instead of doing one-step-back update, we go back further with  $\lambda$  decay

---

**Algorithm 5** SARSA with Neural Network Function Approximation (Explicit Loss Form)

# Markov Decision Processes

- Generally RL courses and books start with MDP formulation of the RL problem
- Remember how the agent's decisions depend only on the current state (e.g.  $a_t \sim \pi_\theta(\cdot | s_t)$ )?
  - That assumption — the current state contains all the relevant information — is called the Markov property
- When this holds, we say the environment is Markov, and that allows us to model it as a Markov Decision Process (MDP)
  - The RL problem is defined as MDP problem:  $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ 
    - $\mathcal{S}$  is the set of all possible states.
    - $\mathcal{A}$  is the set of all the actions available to the agent
    - $P(s' | s, a)$  is **transition probability (how the environment changes)**.
    - $R(s, a)$  is the immediate reward
    - $\gamma$  is the discount factor

# Early RL: Tabular data

- In this lecture we directly went to the use of NNs for modeling the policy functions, value functions etc.
  - In Sutton and Barto, you will see a use of tabular objects:
    - Works if the spaces are limited
    - State-value tables (V-tables)
    - Action-value tables (Q-tables)
    - Transition matrix representing transition probabilities
  - Useful for very very simple problems (like tic-tac-toe) and for understanding algorithms but modern RL usage mostly requires NNs

# Reading only: Dynamic Programming and Bellman Equations

- Transition dynamics are perfectly known as a MDP
  - The algorithms to solve for the optimal policies = Dynamic Programming
  - Read about the bellman equations and proofs of how value iteration converges to optimal policies in Sutton and Barto (Chapter 4 with connections to Chapter 3)
- $$V^\pi(s) = \sum_a \pi(a | s) \sum_{s',r} p(s', r | s, a) [r + \gamma V^\pi(s')]$$
- $$Q^\pi(s, a) = \sum_{s',r} p(s', r | s, a) \left[ r + \gamma \sum_{a'} \pi(a' | s') Q^\pi(s', a') \right]$$

# Off-policy RL

- Both SARSA and REINFORCE that we saw are on-policy
  - We are working with one episode at a time:
    - Can also work with batches for stability similar to supervised RL
  - We can also keep the past experience in a buffer
    - And train the NNs again: **experience replay**
    - However, need to be careful that now you are optimizing a policy using data generated by other: off-policy reinforcement learning (RL)
  - Learning online:
    - Challenging: NNs forget
      - We want them to remember but also not always

# Planning in RL

- We also learn the state-transition functions
  - Think about playing chess again
  - Or any form of imagination:
    - If I do this, this will happen
  - After doing this, you can update your value function for example: this is called planning
  - You have a model of the real world:  
**model-based RL**
- Using tabular algorithm for explanation only
- What else can you do?
- Generative Models
- Used in dreamer architectures
  - <https://arxiv.org/pdf/2301.04104>

---

**Algorithm 6** Dyna-Q

---

**Require:** step size  $\alpha$ , discount factor  $\gamma$ , exploration parameter  $\varepsilon$

**Require:** number of planning steps  $n_{\text{plan}}$

- 1: Initialize action-value estimates  $Q(s, a)$  arbitrarily
- 2: Initialize learned model  $\hat{P}(s' | s, a)$  and  $\hat{R}(s, a)$  (e.g., empty tables)
- 3: **for** episodes  $k = 1, 2, \dots$  **do**
- 4:   Initialize starting state  $s$
- 5:   **for**  $t = 0, 1, 2, \dots$  until  $s$  is terminal **do**
- 6:     Choose action  $a$  using  $\varepsilon$ -greedy policy derived from  $Q(s, \cdot)$
- 7:     Take action  $a$ , observe reward  $r$  and next state  $s'$
- 8:     **(1) Direct RL Update**

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

- 9:     **(2) Update the Model**

Store transition and reward in the model:

$$\hat{P}(s' | s, a) \leftarrow 1, \quad \hat{R}(s, a) \leftarrow r$$

- 10:     **(3) Planning Updates**

**for**  $i = 1, \dots, n_{\text{plan}}$  **do**

Randomly sample a previously observed state-action pair  $(\tilde{s}, \tilde{a})$

Query the learned model to obtain:

$$\tilde{r} = \hat{R}(\tilde{s}, \tilde{a}), \quad \tilde{s}' \sim \hat{P}(\cdot | \tilde{s}, \tilde{a})$$

- 11:       Perform a simulated Q-learning update:

$$Q(\tilde{s}, \tilde{a}) \leftarrow Q(\tilde{s}, \tilde{a}) + \alpha \left[ \tilde{r} + \gamma \max_{a'} Q(\tilde{s}', a') - Q(\tilde{s}, \tilde{a}) \right]$$

- 12:     **end for**


Set  $s \leftarrow s'$

- 13:     **end for**

- 14:     **end for**

# Planning: Monte Carlo Tree Search

- Start with the current state, which becomes the root node of the search tree.
- Selection: From the root, follow a tree policy (e.g., UCB/PUCT) down the tree until you reach a leaf.
- Expansion: If that leaf has any unexplored actions, add one or more new child nodes.
- Simulation: From the leaf (or new child), run a rollout to the end of the episode using a simple policy.
- Backup: Take the return from the rollout and propagate it back up the tree, updating value estimates on all edges used during selection.
- Action choice: After many iterations, pick the action from the root with the highest visit count—this is treated as the best move.



# AlphaGo's Victory (2016)

- First AI to defeat a world champion Lee Sedol in Go on March 15, 2016, a game long considered intractable for classical search.
- Foundation: Monte Carlo Tree Search (MCTS) but with reinforcement learning, and deep neural networks
- Learned from both expert human games and self-play reinforcement learning.
- Used policy networks to propose promising moves and a value network to estimate win probability.



# More successes from MCTS

- AlphaGo Zero (2017) learned only from self-play
- AlphaZero (2017): generalization of AlphaGo Zero to other games
- AlphaFold (2020-2021):
  - Applied deep learning to protein folding, a grand challenge in biology
  - Achieved near-experimental accuracy on protein structure prediction
  - Considered one of the biggest scientific breakthroughs of the decade
- MuZero (2019):
  - General model-based RL algorithm that learns the rules of the environment itself
  - Doesn't need to know game dynamics beforehand

# Using RL for your problem


- Unfortunately, a lot of fine tuning needs to be done. Most algorithms are somewhat unstable
  - A lot of “tricks” (such as gradient clipping) need to be applied to make training works
    - Which is why it is hard to write your own algorithms from scratch: The algorithms we looked at in this lecture are for understanding mostly
    - Play with the parameters to find what works
  - A RL library or your own algorithm
    - RLLib: <https://docs.ray.io/en/latest/rllib/index.html>
    - Stable baselines 3: <https://stable-baselines3.readthedocs.io/en/master/>
    - I always recommend Proximal Policy Optimization (PPO) as the first method to try depending on the problem
- Using a library, we only need two elements. What are they?
  - Defining a value network
  - Defining a policy network

```
.training(  
    train_batch_size=512,  
    minibatch_size=128,  
    num_epochs=4,  
    lr=3e-4,  
    gamma=0.995,  
    lambda_=0.95,  
    use_gae=True,  
    grad_clip=1.0,  
    entropy_coeff=0.005,  
    clip_param=0.8,  
    vf_loss_coeff=0.1,  
    vf_clip_param=1000.0,  
)
```



Training configuration  
for PPO in RLLib

# RL for Chip Placement

- Using PPO
- GNNs as value and policy network
- Achieves competitive or superior PPA (power, performance, area) compared to expert-designed or algorithmic methods, while reducing design iteration time.




# RL for Instrument Design



Sensor type

- Type 3
- Type 2
- Type 1



- <https://doi.org/10.1088/2632-2153/adf7ff>
- Our group works a lot on using AI for design

# Conclusion

- Reinforcement Learning is interesting
  - If you are interested in general theory of learning, look into RL not LLMs
  - Already applied to a lot of problems (chess, design, optimization of many systems)
  - And we are expecting the usage to only increase in the future
- Thank you!